




# Welcome to


## *Introduction to Machine Learning!*




# My research



Zhang et al., 2016




Hu et al., 2012, 2014




Hu et al., 2014



Liang et al., 2013;  
Wu et al., 2013;  
Li et al., 2015





Zhang et al., 2015




BI


AI




Shi et al., 2014



Cheng et al., 2015



Zhang et al., in preparation



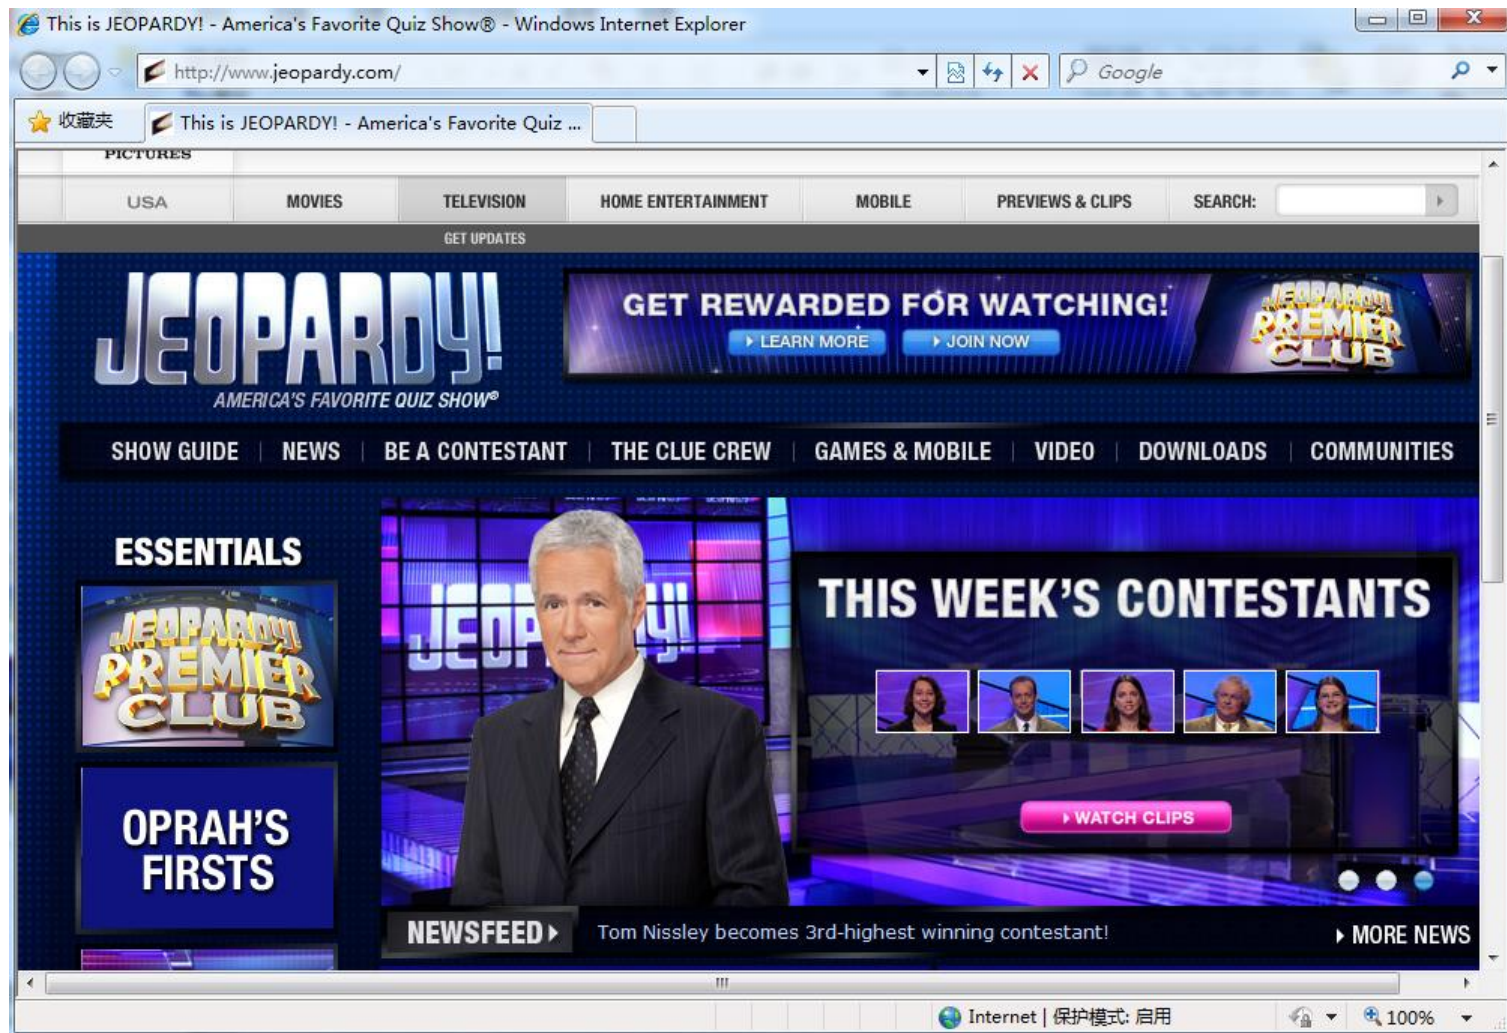
Liang et al., 2015a; 2015b



In progress

# Coffee Time





Jeopardy: An American TV show  
Requires the players to suss out the subtleties of language from  
jokes and puns to irony and anagrams

# IBM Watson @ Jeopardy

- February 14, 15, and 16, 2011
  - *Jeopardy's* two biggest champions
  - Brad Rutter (right):
    - Won a whopping \$3.25 million playing *Jeopardy*, the most cash ever awarded on the show.
    - He is a Johns Hopkins University dropout
  - Ken Jennings (left):
    - Holds the title for longest *Jeopardy* winning streak, with 74 consecutive wins in 2004.
    - He holds degrees in computer science and English, from Brigham Young University, and an international BA diploma from Seoul Foreign School.





# IBM Watson won the Jeopardy


- [http://domino.watson.ibm.com/library/cyberdig.nsf/papers/D12791EAA13BB952852575A1004A055C/\\$File/rc24789.pdf](http://domino.watson.ibm.com/library/cyberdig.nsf/papers/D12791EAA13BB952852575A1004A055C/$File/rc24789.pdf)

Towards the Open Advancement of Question Answering Systems



Final:  
\$77,147  
to  
\$21,600 &  
\$24,000.

# IBM Watson



- In development for 4 years
- Runs on 90 servers
- Does not connect to the Internet
- Search on a large scale knowledge base
- Trained with previous questions and games
  - With Jeopardy players: 77 (2009) + 55 (2010, winners)
  - Category: US Cities
  - *Its largest airport was named for a World War II hero; its second largest, for a World War II battle.*
  - *What is Chicago / Toronto ?*

# Technical requirements

- Answers to questions on any topic
  - Science, geography, popular culture ...
- Accuracy: not only an answer, but a confident right answer
- Speed: within 3 second or less

- Advanced linguistic understanding
  - Parser complex sentences, recognize and understand jokes, metaphors, puns and riddles
- Real time analysis of questions
- Learn from mistakes
- Be prepared to handle the unexpected ...



# Techniques involved -- DeepQA

- A massively parallel probabilistic evidence-based architecture for answering questions
  - Non-database approach
  - Deep text analytics
    - NLP and statistical NLP
  - Machine learning
    - Formulating parallel hypotheses with confidence score
    - Voting, Question interpretation...
  - Search
  - Risk assessment
  - Hadoop and UIMA

# Topic 7. Instance-based learning

Xiaolin Hu

[xlhu@tsinghua.edu.cn](mailto:xlhu@tsinghua.edu.cn)

Updated on Mar 31, 2017

# Motivation


- Previous learning approaches
  - Make an assumption to the model
    - LSE, Decision Tree, MAP, MLE, Naïve Bayes, MM, HMM...
  - Find the optimal parameter

**Is there a learning approach that is NOT  
“model assumption + parameter estimation” ?**




# Motivation

**“Find the 10 most similar images to this image”**



**“Find all closest matching gene segments between two genomes”**



# Some Vocabulary

- Parametric (参数化) vs. Non-parametric (非参数化)
  - Parametric:
    - A particular functional form is assumed
    - Advantage of **simplicity** - easy to estimate and interpret
    - **May have high bias** because the real data may not obey the assumed functional form.
  - Non-parametric:
    - Distribution or density estimate is **data-driven**
    - Relatively few assumptions are made a priori about the functional form.



# Instance-based Learning

- No model is built – Just store all training examples
- Any processing is delayed until a new instance must be classified.


Non-parametric method

# Outline

- Nearest Neighbor
- K-Nearest Neighbor
- KNN Discussion
- Distance-Weighted KNN

# 1. Nearest Neighbor

Similarity  $\leftrightarrow$  distance




# Nearest Neighbor Example

- Credit Rating
  - Classifier
    - Good / Poor
  - Features:
    - L = # late payments/yr;
    - R = Income/Expenses

name	L	R	G/P
A	0	1.2	G
B	25	0.4	P
C	5	0.7	G
D	20	0.8	P
E	30	0.85	P
F	11	1.2	G
G	7	1.15	G
H	15	0.8	P

# Nearest Neighbor Example

name	L	R	G/P
A	0	1.2	G
B	25	0.4	P
C	5	0.7	G
D	20	0.8	P
E	30	0.85	P
F	11	1.2	G
G	7	1.15	G
H	15	0.8	P






# Nearest Neighbor Example (Cont.)

name	L	R	G/P
I	6	1.15	?
J	22	0.45	?
K	15	1.2	?


Distance Measure:

- Scaled distance

$$\sqrt{(L_1 - L_2)^2 + 100(R_1 - R_2)^2}$$



# Nearest neighbor is related to Voronoi diagram



Euclidean distance

- The points  $p_k$  separate the space into adjacent cells  $R_k$
- $R_k$  consists of every point whose distance to  $p_k$  is less than or equal to its distance to any other  $p_k$

From wikipedia

# The problem

- *What if the nearest neighbor is noise?*




- Solution

- Use more than one neighbors
- Voting among neighbors









- **k-Nearest Neighbor**



## 2. K-Nearest Neighbor (KNN)

---

# KNN: example (3-NN)

Customer	Age	Income (K)	#cards	Res	Distance from David
John 	35	35	3	No	$\text{sqrt} [(35-37)^2 + (35-50)^2 + (3-2)^2] = 15.16$
Rachel 	22	50	2	Yes	$\text{sqrt} [(22-37)^2 + (50-50)^2 + (2-2)^2] = 15$
Hannah 	63	200	1	No	$\text{sqrt} [(63-37)^2 + (200-50)^2 + (1-2)^2] = 152.23$
Tom 	59	170	1	No	$\text{sqrt} [(59-37)^2 + (170-50)^2 + (1-2)^2] = 122$
Nellie 	25	40	4	Yes	$\text{sqrt} [(25-37)^2 + (40-50)^2 + (4-2)^2] = 15.74$
David 	37	50	2	Yes	



# KNN discussion

---

# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics

# KNN discussion 1: distance metrics

- Minkowski or  $L_\lambda$  metric:



$$d(i, j) = \left( \sum_{k=1}^p |x_k(i) - x_k(j)|^\lambda \right)^{\frac{1}{\lambda}}$$

- Euclidean Distance ( $\lambda = 2$ )

$$d(i, j) = \sqrt{\sum_{k=1}^p (x_k(i) - x_k(j))^2}$$

- Manhattan dis., city block dis. ( $\lambda = 1$ ) :

$$d(i, j) = \sum_{k=1}^p |x_k(i) - x_k(j)|$$



- Chebyshev distance, chessboard distance,  $L_\infty$


$$d(i, j) = \max_k |x_k(i) - x_k(j)|$$

- Mean Censored Euclidean :


$$\sqrt{\sum_k (x_{ik} - x_{jk})^2 / n}$$

- Bray-Curtis  $\sum_k |x_{ik} - x_{jk}| / \sum_k (x_{ik} + x_{jk})$

- Canberra  $\frac{\sum_k |x_{ik} - x_{jk}|}{\sum_k (x_{ik} + x_{jk})}$

	a	b	c	d	e	f	g	h	
8	5	4	3	2	2	2	2	2	8
7	5	4	3	2	1	1	1	2	7
6	5	4	3	2	1		1	2	6
5	5	4	3	2	1	1	1	2	5
4	5	4	3	2	2	2	2	2	4
3	5	4	3	3	3	3	3	3	3
2	5	4	4	4	4	4	4	4	2
1	5	5	5	5	5	5	5	5	1
	a	b	c	d	e	f	g	h	

# Voronoi diagram with different distance metrics



Euclidean distance



Manhattan distance


From wikipedia



# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes

# KNN discussion 2: attributes



**John:**

Age=35

Income=95K

No. of credit cards=3



**Rachel:**

Age=41







Income=215K

No. of credit cards=2

$$\text{Dis (John, Rachel)} = \sqrt{[(35-41)^2 + (95,000-215,000)^2 + (3-2)^2]}$$

- Distance between neighbors could be dominated by some attributes with relatively large numbers
  - e. g. Income
- Important to normalize some features
  - e. g., map numbers to [0-1]

# KNN: attributes normalization

Customer	Age	Income (K)	#cards	Response
John 	$35/63 = 0.55$	$35/200 = 0.175$	$3/4 = 0.75$	No
Rachel 	$22/63 = 0.34$	$50/200 = 0.25$	$2/4 = 0.5$	Yes
Hannah 	$63/63 = 1$	$200/200 = 1$	$1/4 = 0.25$	No
Tom 	$59/63 = 0.93$	$170/200 = 0.85$	$1/4 = 0.25$	No
Nellie 	$25/63 = 0.39$	$40/200 = 0.2$	$4/4 = 1$	Yes
David 	$37/63 = 0.58$	$50/200 = 0.25$	$2/4 = 0.5$	Yes

# KNN: weighted attributes

- The classification of an example is based on all the attributes
  - Independent of their relevance – Even the irrelevant attributes are used.

- Weight the contribution of each attribute based on its relevance, e.g.

$$d_{WE}(i, j) = \left( \sum_{k=1}^p w_k (x_k(i) - x_k(j))^2 \right)^{\frac{1}{2}}$$

- Scaling of dimensions in the distance space
  - $w_k = 0 \rightarrow$  eliminate the corresponding dimension (feature selection)
- Possible weighting method:
  - use mutual information  $I(\text{each attribute}, \text{the class})$
  - Determine automatically by cross validation

# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes
    - Normalization, Weighting
  - Continuous valued target function

# KNN discussion 3:


## continuous valued target function

- Discrete output -- voting
- Continuous valued target function
  - Mean value of the  $k$  nearest training examples


# Continuous valued target function

Red: known instances  
Blue: estimation


## 1-nearest neighbor



## 3-nearest neighbor



## 5-nearest neighbor



# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes
    - Normalization, Weighting
  - Continuous valued target function
  - Choose  $k$



# KNN discussion 4: choose $k$


- In many cases  $k=3$
- Depending on the number of training examples
  - Larger  $k$  doesn't mean better performance
- Cross validation
  - Leave-one-out (Throw-one-out, Hold-one-out)
    - Each time: Take one sample to test, the others are all training examples
- KNN is stable
  - Small perturbation of training data does not change results significantly.

# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes
    - Normalization, Weighting
  - Continuous valued target function
  - Choose k
  - Break ties

# KNN discussion 5: break ties

- What if  $k=3$ , and each near neighbor belongs to a different class?
  - $P(w | X) = 1/3$
  - or find a new neighbor (4<sup>th</sup>)
  - or choose the nearest one
  - or random select one
  - or ...




# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes
    - Normalization, Weighting
  - Continuous valued target function
  - Choose k
  - Break ties
  - More on efficiency

# KNN discussion 6: more on efficiency

- We can speed up the search for the nearest neighbor:
  - Examine nearby points first.
  - Ignore any points that are further than the nearest point found so far.
- Do this using a **KD-tree**:
  - KD-tree: k dimensional tree (the dimension of the points is k)
  - Tree-based data structure
  - Recursively partitions points into axis aligned boxes.

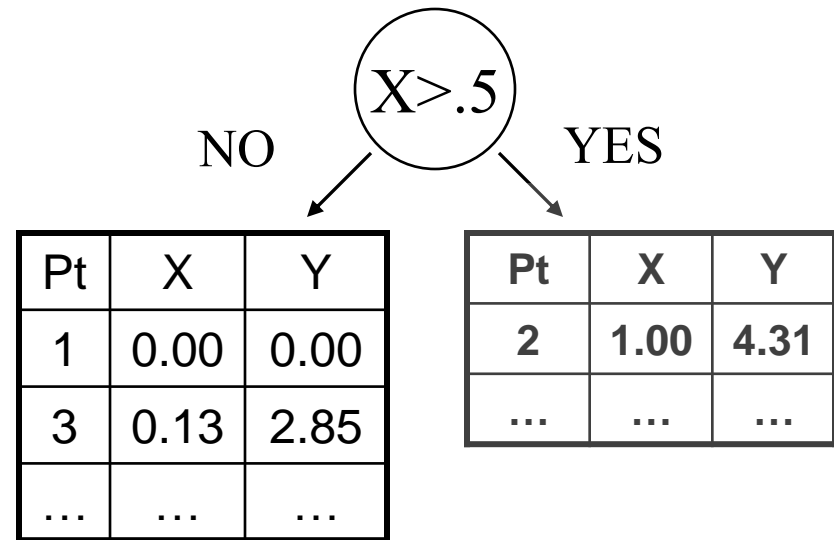

# KD-Tree: Construction



Pt	X	Y
1	0.00	0.00
2	1.00	4.31
3	0.13	2.85
...	...	...

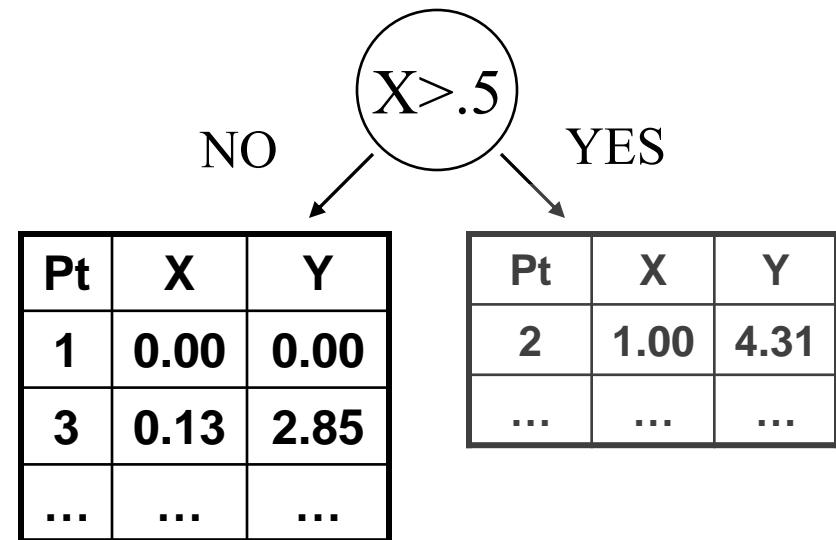

We start with a list of points.

# KD-Tree: (1) Construction



We can split the points into 2 groups by choosing a dimension  $X$  and value  $V$  and separating the points into  $X > V$  and  $X \leq V$ .

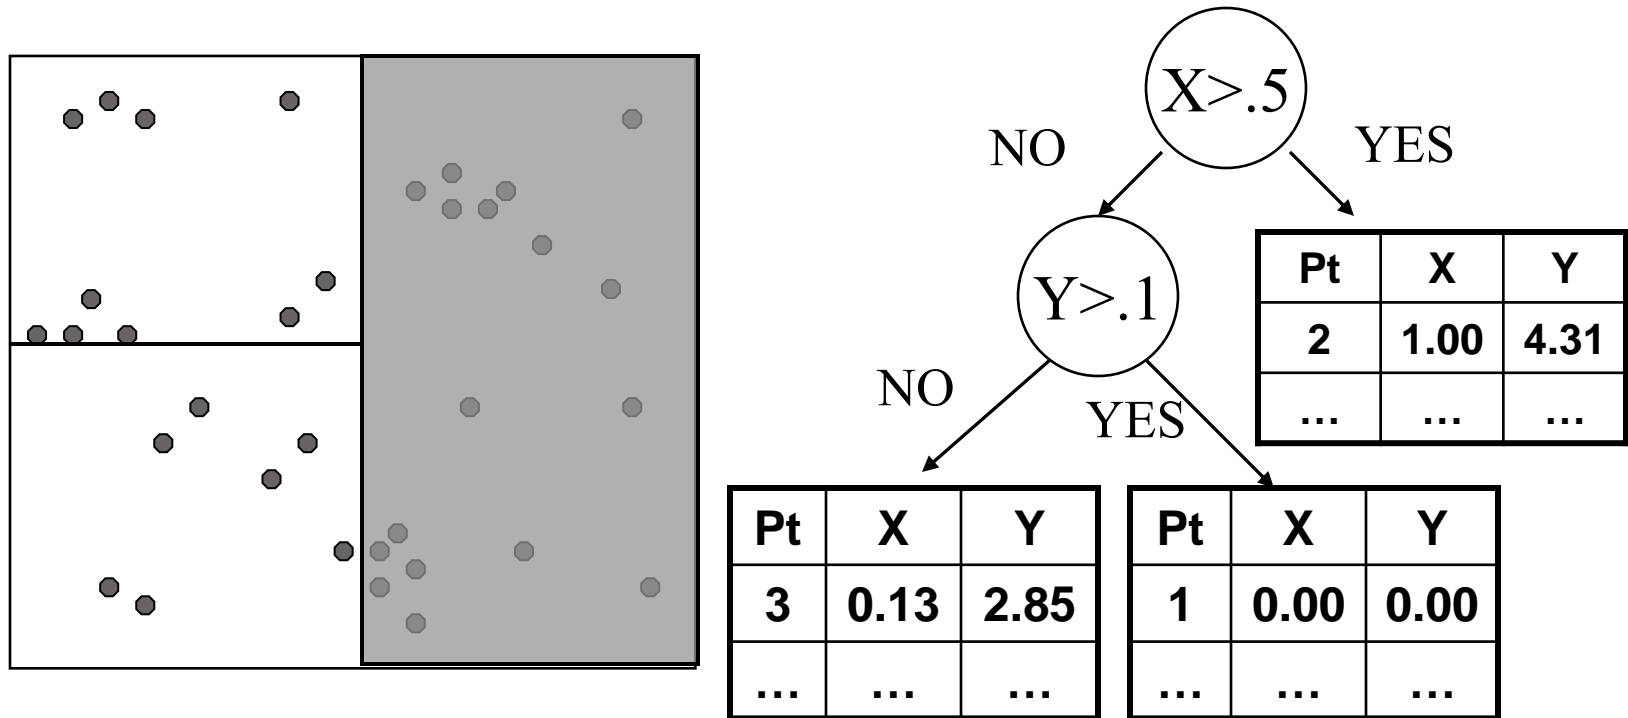
# KD-Tree: (1) Construction



We can then consider each group separately and possibly split again (along same/different dimension).





# KD-Tree: (1) Construction




We can then consider each group separately and possibly split again ([along same/different dimension](#)).

# KD-Tree: (1) Construction



We can keep splitting the points in each set to create a tree structure. Each leaf node contains a list of points.

# KD-Tree: (1) Construction





We will keep around one additional piece of information at each node: **The (tight) bounds of the points at or below this node.**

# KD-Tree: (1) Construction

Use heuristics to make splitting decisions:



- Which dimension do we split along?
  - Widest
- Which value do we split at?
  - Median of value of that split dimension for the points.
- When do we stop?
  - When there are fewer than  $m$  points left, OR
  - The box has hit some minimum width.

## KD-Tree: (2) Query





We traverse the tree looking for the nearest neighbor of the query point.

## KD-Tree: (2) Query





Examine nearby points first: Explore the branch of the tree that is closest to the query point first.

## KD-Tree: (2) Query



Examine nearby points first: Explore the branch of the tree that is closest to the query point first.



## KD-Tree: (2) Query



When we reach a leaf node: compute the distance to each point in the node.





## KD-Tree: (2) Query





When we reach a leaf node: compute the distance to each point in the node.

## KD-Tree: (2) Query





Then we can backtrack and try the other branch at each node visited.

## KD-Tree: (2) Query





Each time a new closest node is found, we can update the distance bounds.

## KD-Tree: (2) Query





Using the distance bounds and the bounds of the data below each node, we can prune parts of the tree that could NOT include the nearest neighbor.

## KD-Tree: (2) Query



Using the distance bounds and the bounds of the data below each node, we can prune parts of the tree that could NOT include the nearest neighbor.

## KD-Tree: (2) Query



Using the distance bounds and the bounds of the data below each node, we can prune parts of the tree that could NOT include the nearest neighbor.

# KNN overview

- Basic algorithm
- Discussion
  - More distance metrics
  - Attributes
    - Normalization, Weighting
  - Continuous valued target function
  - Choose k
  - Break ties
  - More on efficiency - k-Dtree
- Strength and weakness

# KNN overview: strength

- Conceptually simple, yet able to form complex decision boundaries (and complex target functions)
  - e.g. image classification
- Robust to noisy data by averaging k-nearest neighbors
- Comprehensible
  - Easy to explain prediction (near neighbor)
- Information present in the training examples is never lost
  - Because the examples themselves are stored explicitly.
- Simple to implement, stable, no parameters (except for k), easy to run leave-one-out test.





# KNN overview: weakness

- Memory cost
  - Need a lot of space to store all examples
  - Generally, it requires distances between all pairs of points  $O(n^2)$ 
    - K-DTrees  $O(n \log n)$
- CPU cost
  - Takes more time to classify a new example (therefore in offline application)
- It is difficult to determine an appropriate distance function
  - Especially when examples are represented as complex symbolic expressions.
- Irrelevant features have a negative impact on the distance metric.

# The next problem

- Recall: Use more than one neighbors to be robust to noise data

*Do these neighbors have same contributions?*



- Solution
  - Weighting the data
  - Give larger weights to neighbors closer to the query



**Distance-weighted Nearest Neighbor**


### 3. Distance weighted KNN

---


# Distance Weighted KNN

- A weighting function
  - $w_i = K(d(x_i, x_q))$
  - $d(x_i, x_q)$  is the distance between query and  $x_i$
  - $K$  - kernel function that determines the weight for each point
- Output:
  - weighted average:  $\text{predict} = \Sigma w_i y_i / \Sigma w_i$
- Kernel function  $K(d(x_i, x_q))$ 
  - $1/d^2, e^{-d}, 1/(1+d), \dots$
  - Should vary **inversely with the distance  $d$**

# Recall




# Recall




# Distance Weighted NN

$$W_i(d(x_q, x_i)) = 1/d(x_q, x_i)^2$$



# Distance Weighted NN


$$W_i(d(x_q, x_i)) = 1/(d_0 + d(x_q, x_i))^2$$





# Distance Weighted NN

$$W_i(d(x_q, x_i)) = e^{-(d(x_q, x_i)/\sigma_0)^2}$$



Overview:

A memory based learner - 4 factors

---

# A memory based learner: 4 factors

1. A distance metric
2. How many nearby neighbors to look at?
3. A weighting function (optional)
4. How to fit with the local points?

# 1-NN

A memory based learner: 4 factors

1. A distance metric

Euclidian


2. How many nearby neighbors  
to look at? One

3. A weighting function (optional)

Unused

4. How to fit with the local  
points?

the same as the nearest neighbor.



# K-NN

A memory based learner: 4 factors

1. A distance metric

Euclidian


2. How many nearby neighbors  
to look at?  $K$

3. A weighting function (optional)

Unused

4. How to fit with the local  
points?

Voting among  $K$  neighbors



# Distance-weighted KNN

A memory based learner: 4 factors

1. A distance metric

Euclidian

2. How many nearby neighbors

to look at? All of them, or K

3. A weighting function (optional)

*e. g.*  $w_i = \exp(-D(x_i, \text{query})^2 / K_w^2)$


$K_w$ : Kernel Width. Very important

4. How to fit with the local points:

the weighted average of the outputs

$\text{predict} = \Sigma w_i y_i / \Sigma w_i$

$$W_i(d(x_q, x_i)) = e^{-(d(x_q, x_i)/\sigma_0)^2}$$

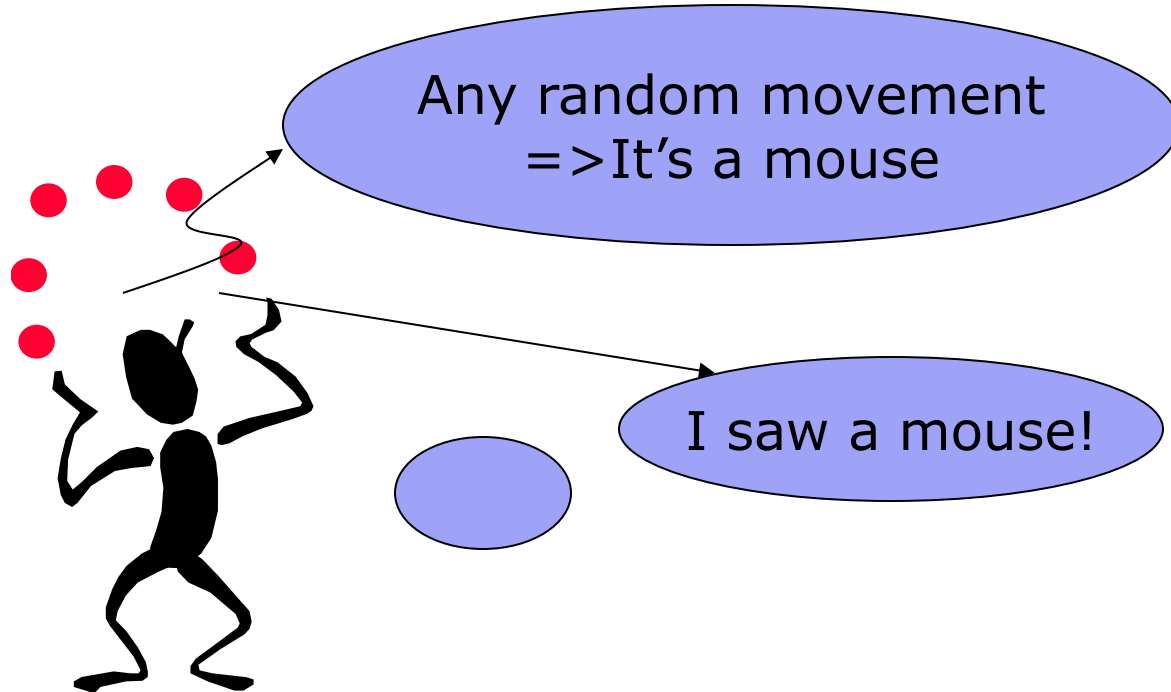


# Lazy and eager learning

---

# Different learning methods

- Eager Learning





# Different learning methods

- Instance-based Learning (lazy learning)



Its very similar to a Desktop!!



# Lazy vs. Eager Learning

- Lazy: wait for query before generalizing
  - Training time: short
  - Test time: time consuming
- Lazy learner:
  - Can create local approximations
- Eager: generalize before seeing query
  - Training time: long
  - Test time: short
- Eager learner:
  - Use same model to each query
  - Tend to create global approximation

If they use the same hypothesis space, lazy can represent more complex functions